




uSAP: An Ultra-Fast Stochastic Graph Partitioner

2023 MIT/IEEE/Amazon HPEC Graph Challenge

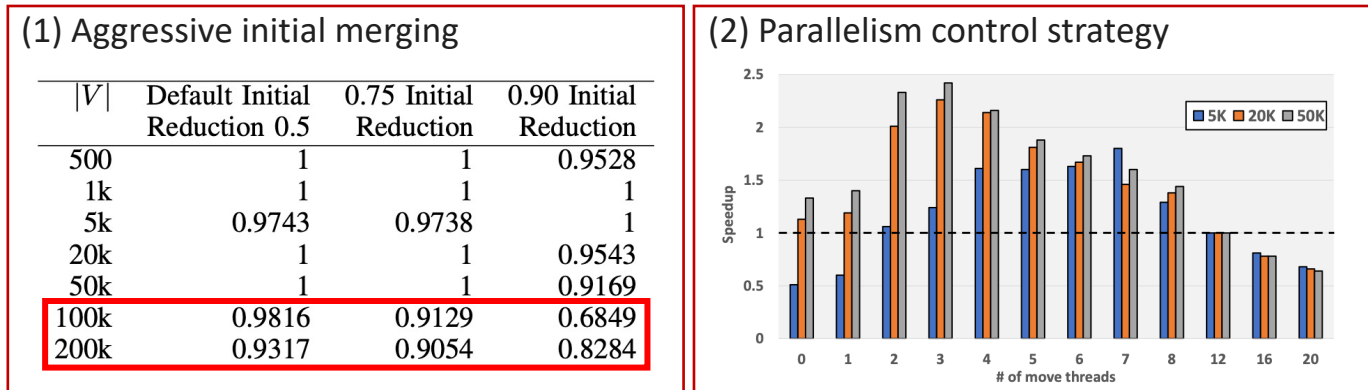


Challenges

- **Time-consuming** baseline sequential partitioner (PEIXOTO)



- **Improvable results** from 2021 Champion (FSBP)



Recall decrease!


Not scalable!



Solutions

- **Strongly Connected Components**-based initial block merging
⇒ tries to find the blocks with more interactions to merge


(1) Merge the block B



$$M_1^+ = \begin{pmatrix} A & B & C & D & E \\ 1 & 0 & 2 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 \end{pmatrix} \begin{matrix} A \\ B \\ C \\ D \\ E \end{matrix}$$

$$\Rightarrow \Delta E = 0.35$$

(2) Merge the block C



$$M_2^+ = \begin{pmatrix} A & B & C & D & E \\ 2 & 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 \end{pmatrix} \begin{matrix} A \\ B \\ C \\ D \\ E \end{matrix}$$

$$\Rightarrow \Delta E = -0.43$$


We use a tunable threshold t_{SCC} to determine when to stop.




Solutions

- **Batch Parallel Nodal Block Assignment**


(1) Randomly select a batch



(2) Nodal block assignment in parallel



(3) Update the shared state

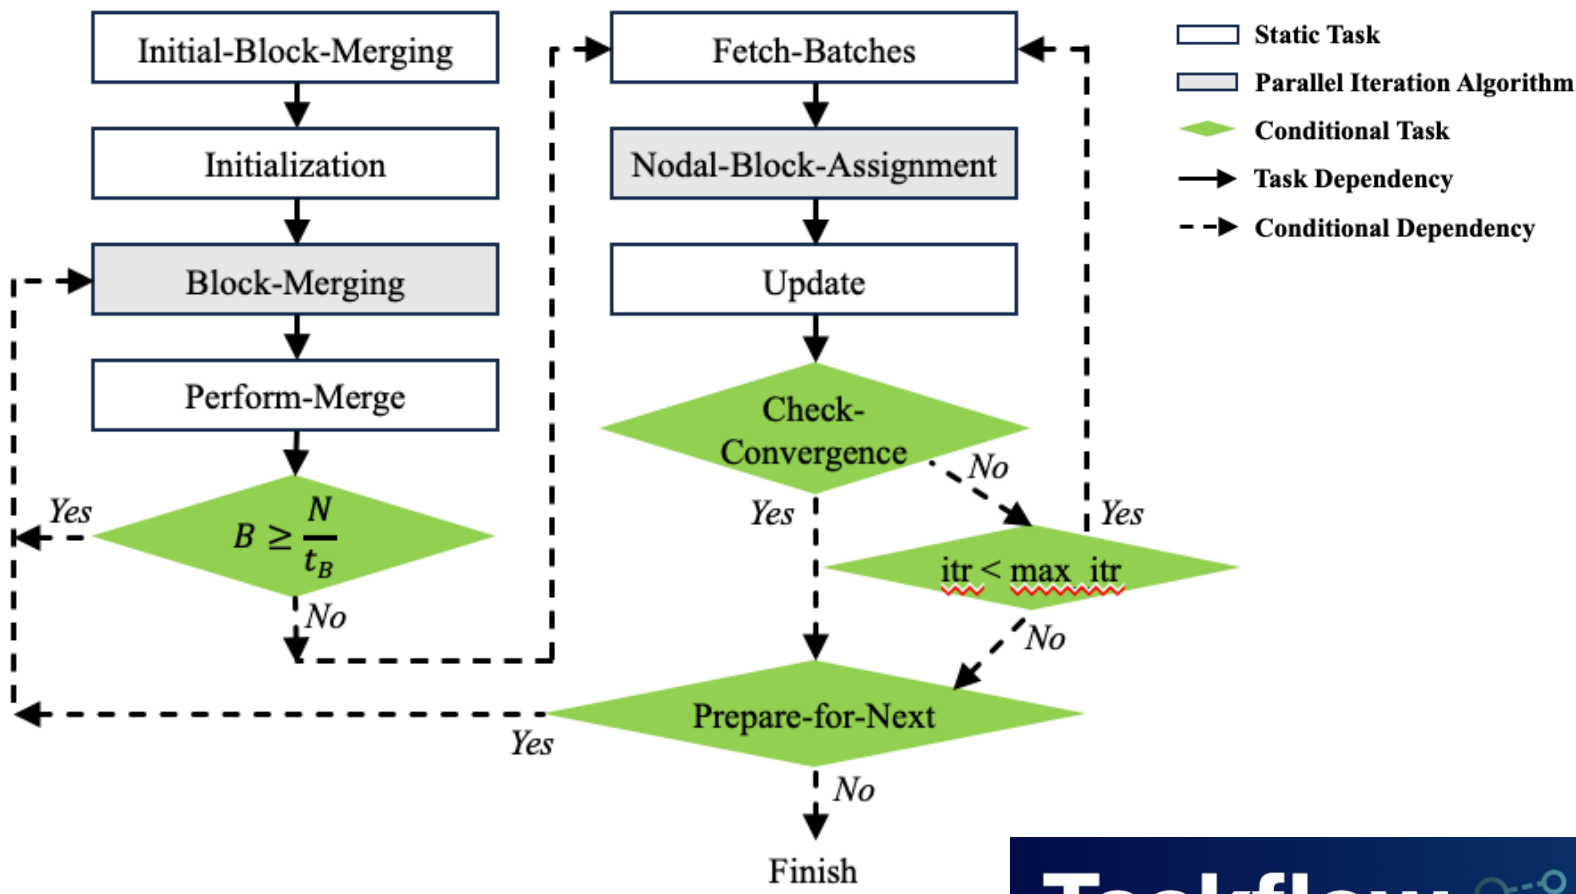


The results of (2) are saved in a shared vector.



Solutions

- Task Graph Parallelism



Whether to perform nodal block assignment or not



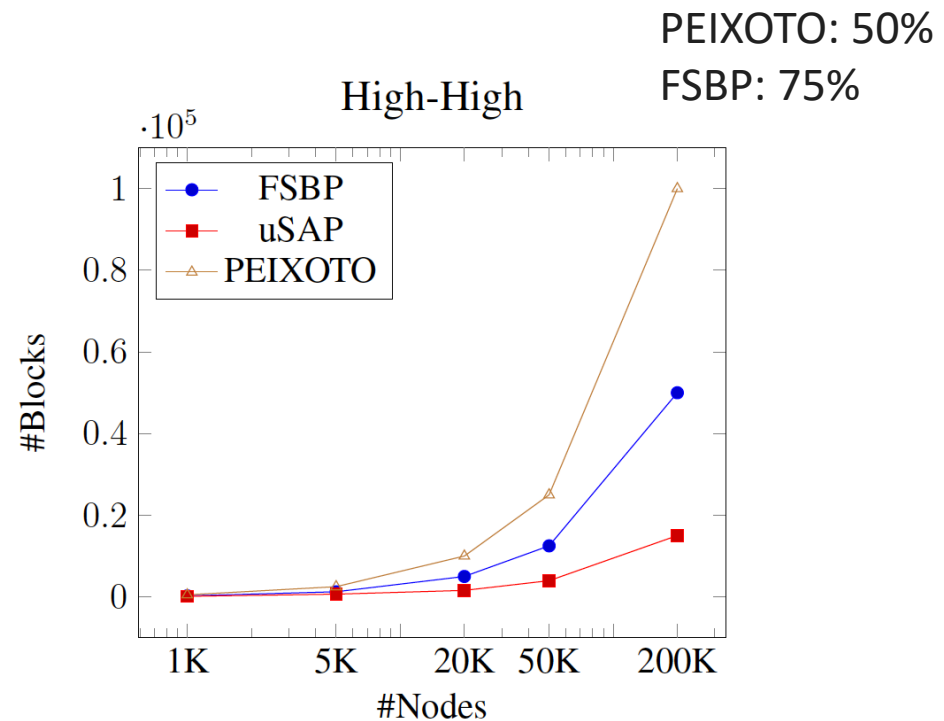

Experimental Setup

- **Baseline**
 - Sample code provided by Graph Challenge (PEIXOTO)
 - Faster Stochastic Block Partition (FSBP)
- **Dataset**
 - 2022 Streaming Graph Partition Dataset provided by Graph Challenge
- **Software**
 - Ubuntu Linux 5.15.0-58-generic x86_64 machine
 - GNU GCC-11.3.0 with C++17
- **Hardware**
 - 12-core Intel® Core™ i7-12700 processor
 - 32GB RAM



Experimental Result

- Number of blocks after the initial block merging





Overall Runtime

Static Graph Categories												
Nodes	Low-Low			Low-High			High-Low			High-High		
	PEIXOTO	FSBP	uSAP	PEIXOTO	FSBP	uSAP	PEIXOTO	FSBP	uSAP	PEIXOTO	FSBP	uSAP
	sec MB	sec MB	sec MB	sec MB	sec MB	sec MB	sec MB	sec MB	sec MB	sec MB	sec MB	sec MB
1K	47.3	8.5	0.005	46.8	8.6	0.006	59.9	9.9	0.007	58.9	10.1	0.007
	86.6	49.1	6.9	86.4	48.0	7.3	86.7	49.3	7.4	86.6	48.8	7.2
5K	470.0	56.2	0.7	522.8	70.5	0.6	596.5	75.8	1.2	550.7	71.3	1.1
	244.6	78.9	13.6	244.7	78.5	15.7	249.9	78.6	14.3	239.1	79.3	16.7
20K	5305.5	640.6	6.2	5450.9	641.2	7.0	5681.6	689.2	10.4	5176.6	607.7	10.1
	2074.4	378.2	126.1	1951.7	386.4	136.1	2073.3	385.3	136.1	2033.6	382.0	132.6
50K	>36000	3558.6	27.5	>36000	3462.6	27.4	>36000	3381.4	53.5	>36000	3419.2	41.4
	-	942.6	318.9	-	981.5	310.5	-	947.3	310.5	-	970.8	322.5
200K	>36000	>36000	299.8	>36000	>36000	338.8	>36000	>36000	1243.8	>36000	>36000	1381.2
	-	-	1222.6	-	-	1228.8	-	-	1228.8	-	-	1334.5

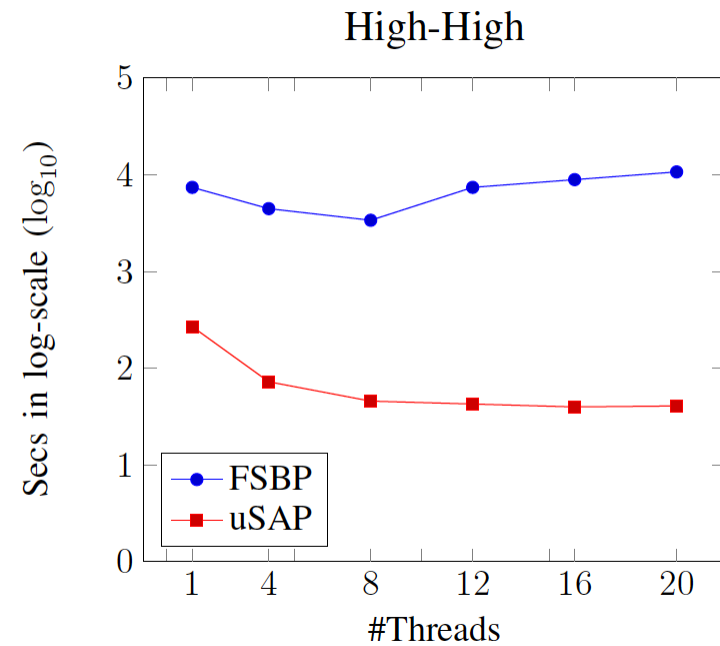



Pairwise Precision and Recall

Static Graph Categories												
Nodes	Low-Low			Low-High			High-Low			High-High		
	PEIXOTO	FSBP	uSAP	PEIXOTO	FSBP	uSAP	PEIXOTO	FSBP	uSAP	PEIXOTO	FSBP	uSAP
	PP PR	PP PR	PP PR	PP PR	PP PR	PP PR	PP PR	PP PR	PP PR	PP PR	PP PR	PP PR
1K	0.994	0.994	0.998	0.939	0.811	0.952	0.717	0.719	0.747	0.686	0.710	0.843
	0.996	0.996	0.998	0.990	0.995	0.993	0.883	0.878	0.971	0.972	0.653	0.954
5K	0.940	1.000	1.000	0.976	0.970	0.987	0.861	0.641	0.983	0.676	0.689	0.861
	1.000	1.000	1.000	0.769	0.850	0.998	0.816	0.666	0.997	0.789	0.721	0.801
20K	0.984	1.000	1.000	0.721	0.921	0.948	0.950	0.875	0.982	0.889	0.789	0.803
	1.000	1.000	1.000	0.671	0.479	0.999	0.740	0.705	0.999	0.995	0.943	0.778
50K	-	0.988	1.000	-	0.849	0.921	-	0.757	0.946	-	0.766	0.456
	-	0.855	1.000	-	0.997	0.998	-	0.862	0.994	-	0.64	0.981
200K	-	-	0.983	-	-	0.768	-	-	0.832	-	-	0.386
	-	-	1.000	-	-	0.922	-	-	0.314	-	-	0.885



Scalability





Contribution

- SCC-based Initial block merging
- Dynamic batch parallel nodal block assignment
- Task graph parallelism



Thank you!

—
uSAP: An Ultra-Fast Stochastic Graph Partitioner

<https://github.com/gary30404/uSAP.git>